

Tutorial on RDF Stream Processing

M. Balduini, J-P Calbimonte, O. Corcho,
D. Dell'Aglio, E. Della Valle

<http://streamreasoning.org/rsp2014>



SPARQLStream: Ontology- based access to data streams

Jean-Paul Calbimonte, Oscar Corcho

jp.calbimonte@upm.es, ocorcho@fi.upm.es

<http://www.oeg-upm.net/>



- This work is licensed under the Creative Commons Attribution 3.0 Unported License.

- **You are free:**



to Share — to copy, distribute and transmit the work



to Remix — to adapt the work

- **Under the following conditions**



Attribution — You must attribute the work by inserting

- “[source <http://streamreasoning.org/sr4ld2013>]” at the end of each reused slide
- a credits slide stating
 - These slides are partially based on “Streaming Reasoning for Linked Data 2013” by M. Balduini, J-P Calbimonte, O. Corcho, D. Dell'Aglio, E. Della Valle, and J.Z. Pan <http://streamreasoning.org/sr4ld2013>

- To view a copy of this license, visit <http://creativecommons.org/licenses/by/3.0/>

- Virtual RDF views over data streams
- Ontology-based access to data streams
 - Examples
 - Architecture
 - Underlying query processors
- SPARQLStream language
- Query rewriting
 - R2RML mappings
- Resources

Querying RDF Streams

users, applications



W3C SPARQL

+streams



RDF Stream

W3C RDF

+streams

query processing




RDF Stream Processor

Where is the **data** coming from?



Existing streaming
data sources

DSMS



CEP



Sensor
middleware



...

Wrap
Import
Load
Continuously

Virtual RDF views over data streams

users, applications



W3C SPARQL

Stream



Virtual RDF Stream

W3C RDF

query processing




RDF Stream Processor

Morph-streams

queries



DSMS



CEP



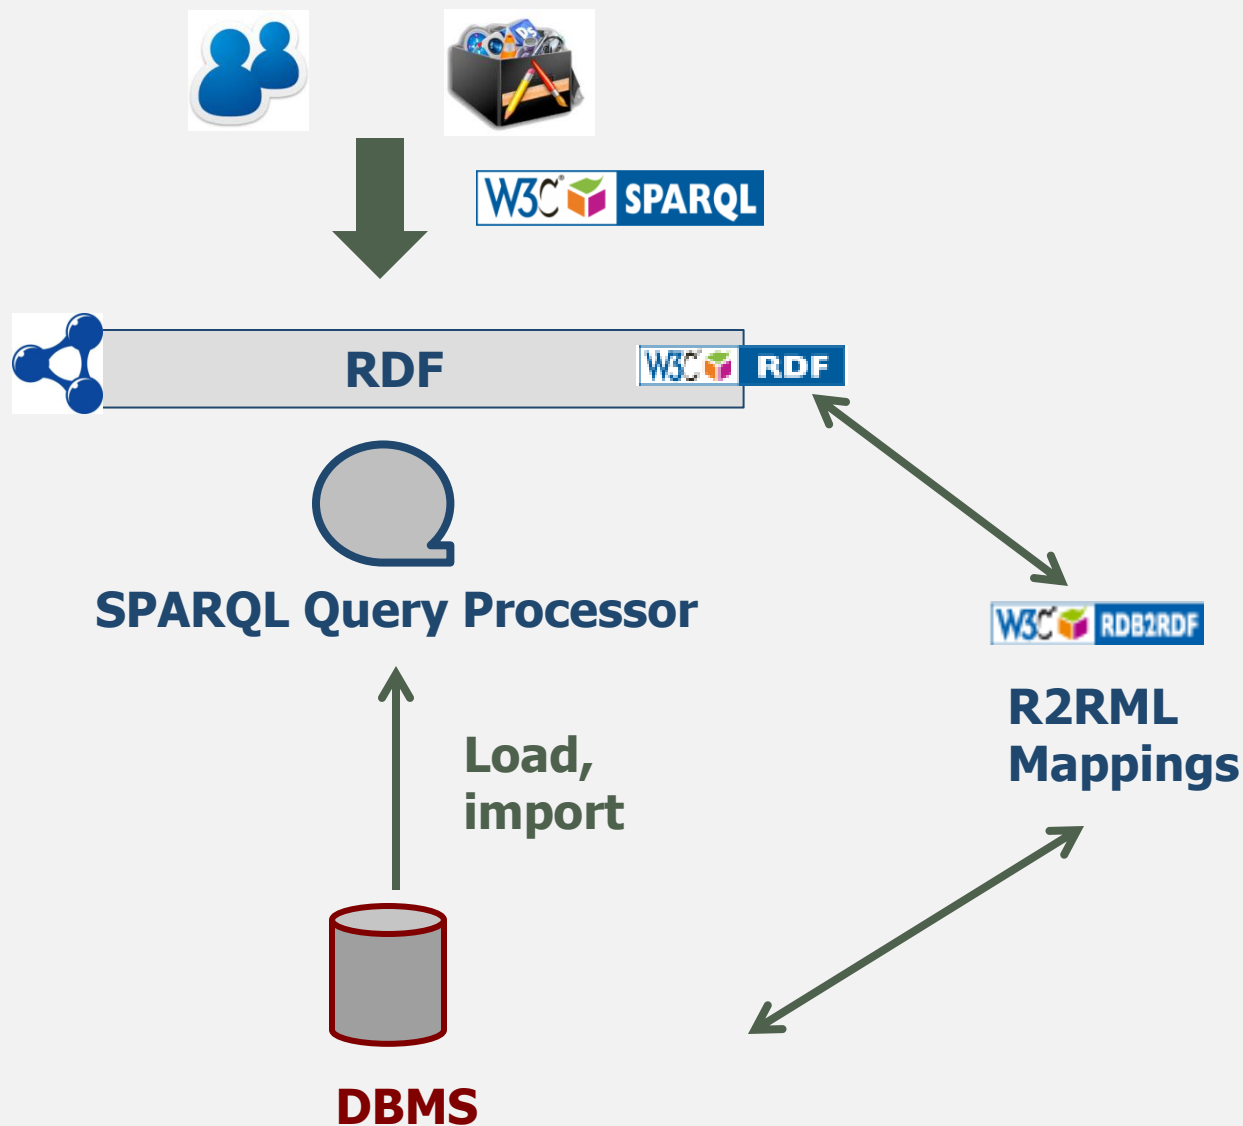
Sensor
middleware



...

data layer

Already seen somewhere...?



Stream Processor Implementations

Data Stream Management Systems (DSMS)

CQL/Stream Borealis
StreamMill NiagaraCQ
TelegraphCQ Esper

Complex Event Processors (CEP)

GEM Rapide Oracle CEP
Cayuga CEDR IBM InfoSphere

Stream Data Middleware

Hourglass Cosm Microsoft StreamInsight
SStreamWare GSN Sybase CEP StreamBase

Diverse **query languages**

Different **query capabilities**

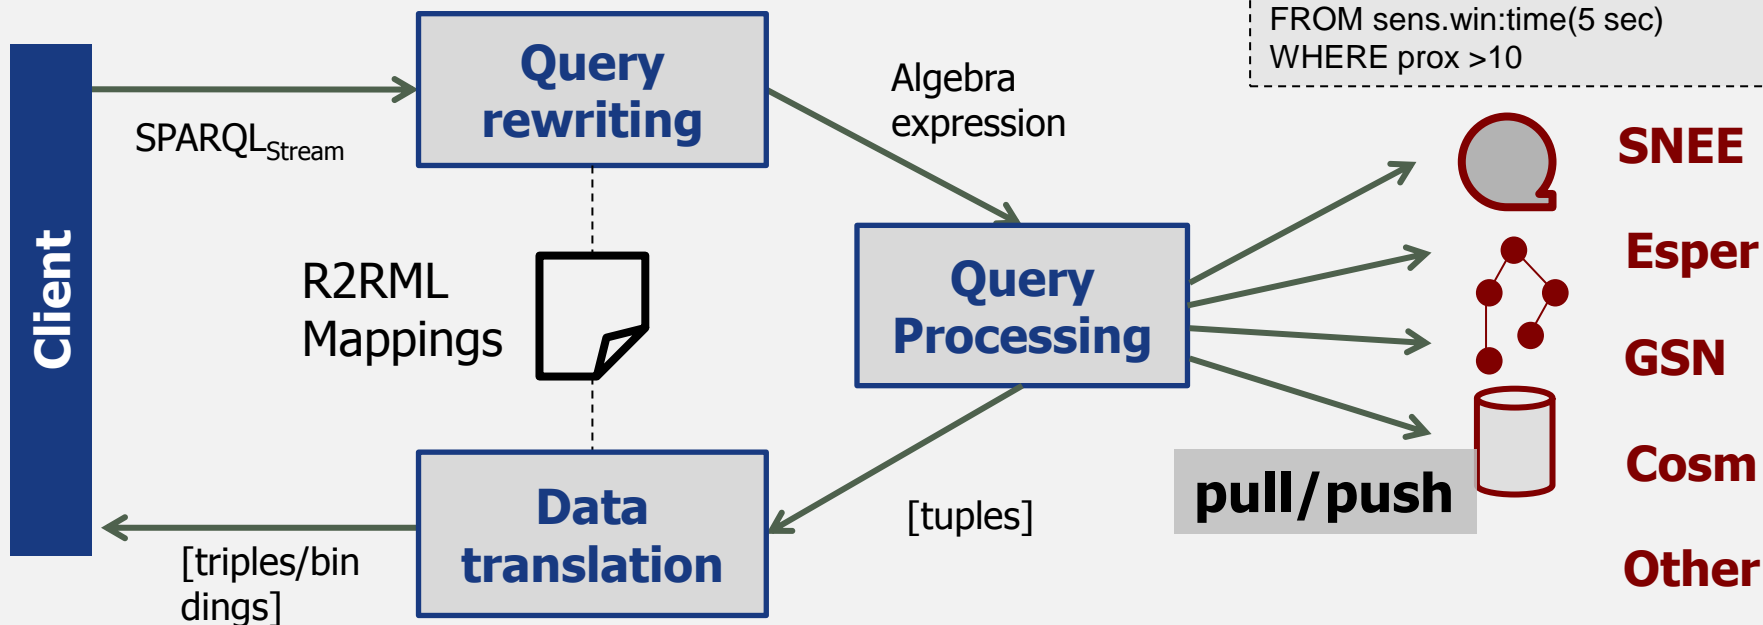
Different **query models**

Morph-streams: Overview

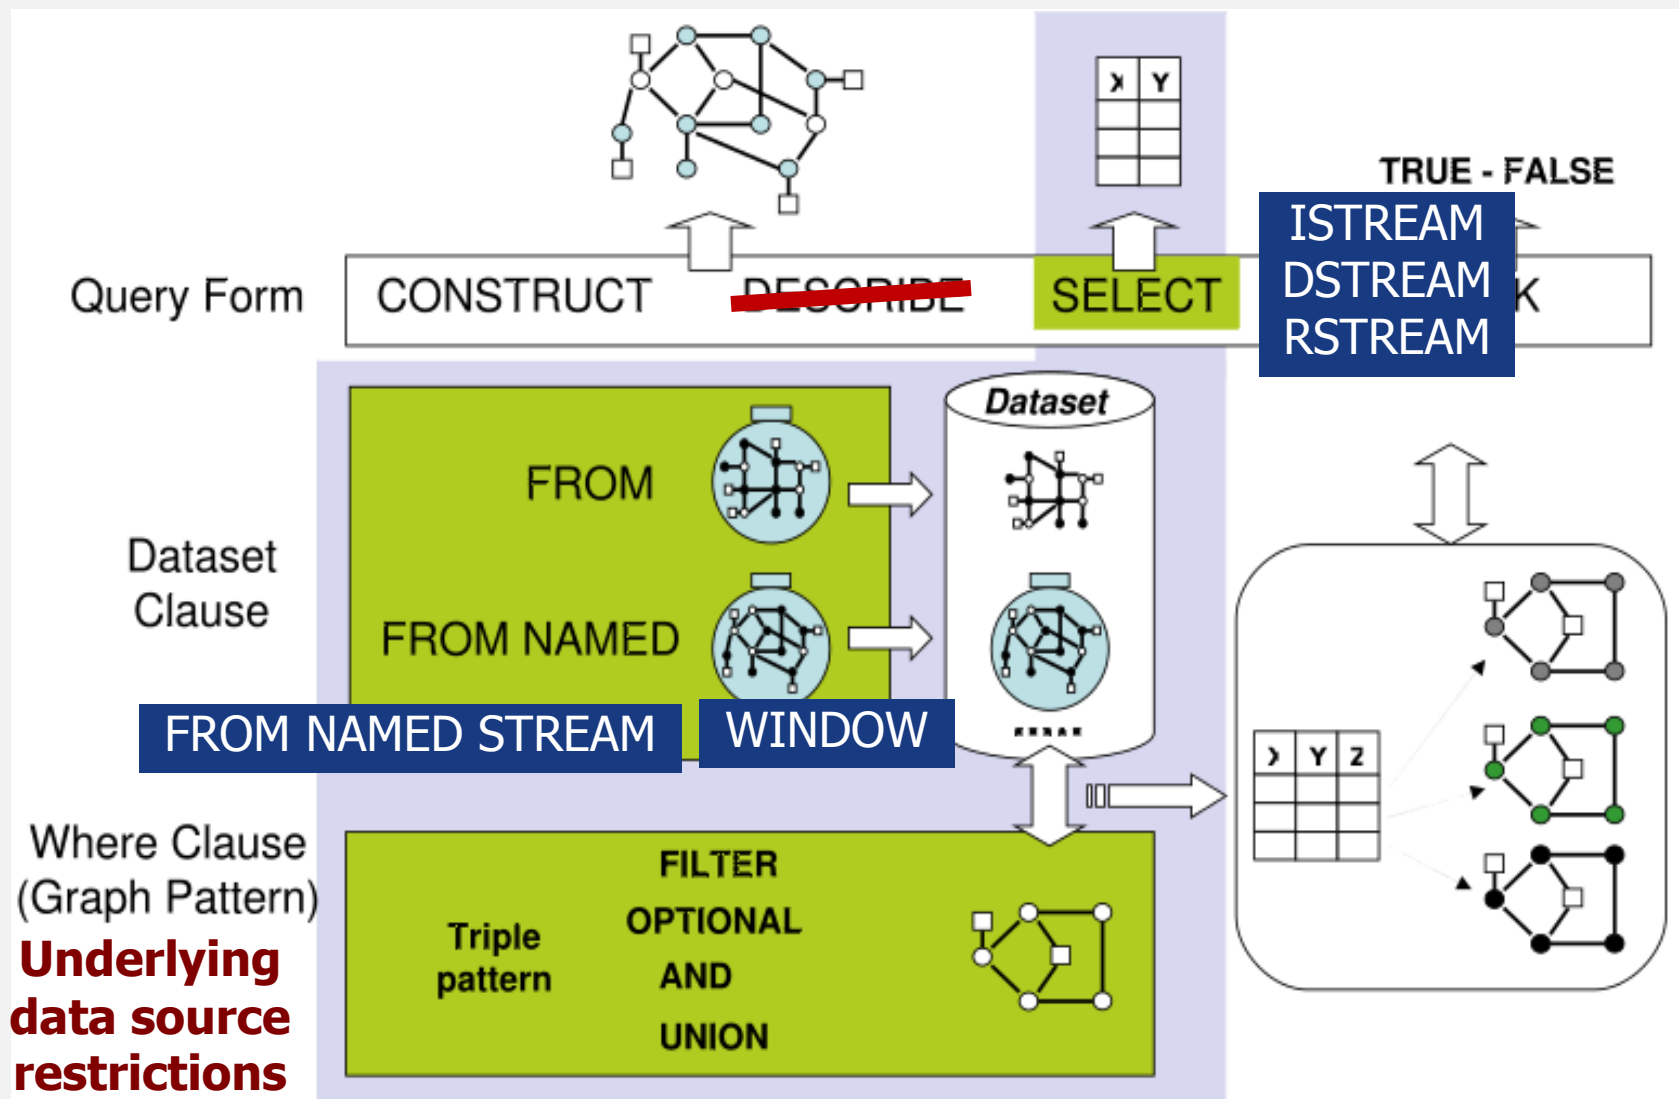
```
SELECT ?proximity
FROM STREAM
<http://streamreasoning.org/SensorReadings.srdf> [NOW-5
S]
WHERE {
  ?obs a ssn:ObservationValue;
  qudt:numericalValue ?proximity;
  FILTER (?proximity>10) }
```

$\Pi_{\text{timed,prox}}$
 $\sigma_{\text{prox}>10}$
 $\omega_{5 \text{ Seconds}}$
sens

```
SELECT prox
FROM sens.win:time(5 sec)
WHERE prox >10
```



Morph-streams processing SPARQL_{Stream} queries



SPARQL_{Stream}

```
PREFIX sr4ld: <http://www.streamreasoning.org/ontologies/socialsensor,owl#>
SELECT ?room
FROM NAMED STREAM <http://www.streamreasoning.org/streams/socialsensor.srdf> [NOW-10 S]
WHERE {
    ?obs sr4ld:observedBy ?sensor.
    ?obs sr4ld:where ?room.
}
```

All rooms where something was observed in the last 10s

```
PREFIX sr4ld: <http://www.streamreasoning.org/ontologies/socialsensor,owl#>
SELECT (COUNT(?person) AS ?nmb) ?room
FROM NAMED STREAM <http://www.streamreasoning.org/streams/socialsensor.srdf> [NOW-10 S]
WHERE {
    ?obs sr4ld:who ?pers.
    ?obs sr4ld:where ?room.
}
GROUP BY ?room
```

Number of persons observed in each room in the last 10s

- *NamedStream* → 'FROM' ['NAMED'] 'STREAM' *StreamIRI* '[' *Window* '']
- *Window* → 'NOW-' *Integer TimeUnit* [*UpperBound*] [*Slide*]
- *UpperBound* → 'TO NOW-' *Integer TimeUnit*
- *Slide* → 'SLIDE' *Integer TimeUnit*
- *TimeUnit* → 'MS' | 'S' | 'MINUTES' | 'HOURS' | 'DAY'

```
SELECT ISTREAM ?room
FROM NAMED STREAM <http://www.streamreasoning.org/streams/socialsensor.srdf> [NOW-10 S]
WHERE {...
```

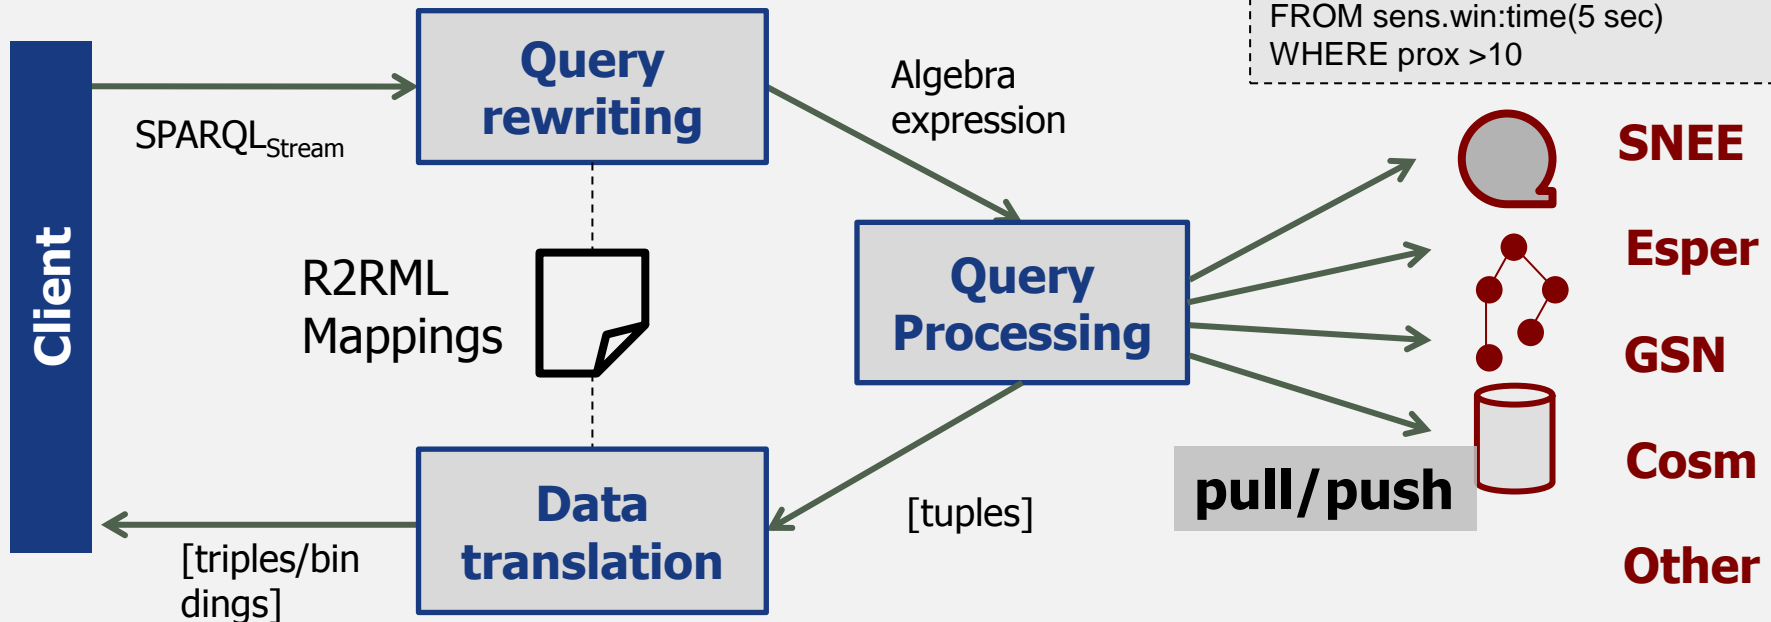
- *Select* → 'SELECT' [*Xstream*] [*Distinct* | *Reduced*] ...
- *Xstream* → 'RSTREAM' | 'ISTREAM' | 'DSTREAM'

Morph-streams: Overview

```
SELECT ?proximity
FROM STREAM
<http://streamreasoning.org/SensorReadings.srdf> [NOW-5
S]
WHERE {
  ?obs a ssn:ObservationValue;
  qudt:numericalValue ?proximity;
  FILTER (?proximity>10) }
```

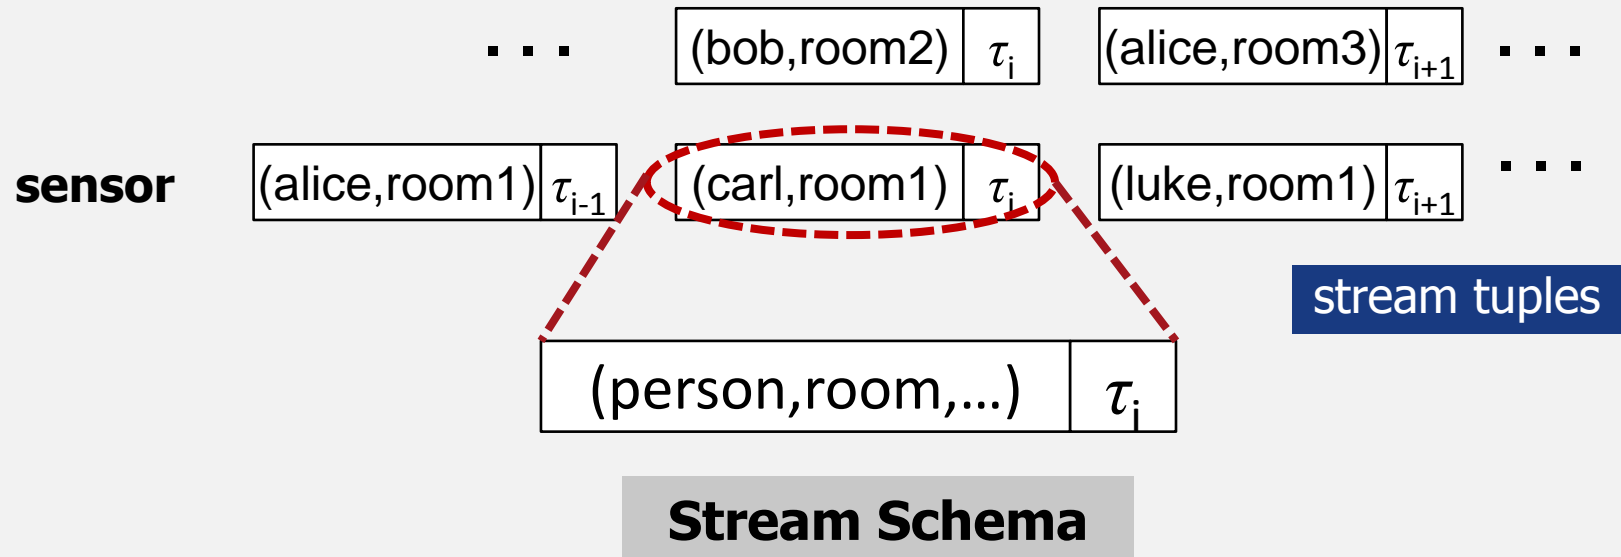
$\Pi_{\text{timed,prox}}$
 $\sigma_{\text{prox}>10}$
 $\omega_{5 \text{ Seconds}}$
sens

```
SELECT prox
FROM sens.win:time(5 sec)
WHERE prox >10
```



Morph-streams processing SPARQL_{Stream} queries

Now, where is the data?



DSMS, CEP, middleware can
evaluate queries over this model

Esper

- CEP/DSMS
- EPL language

```
SELECT prox FROM sensors.win:time(5 minute)
WHERE prox >10
```

SNEE

- DSMS/Sensor Network Query Evaluator
- Compile queries to sensor code

```
SELECT prox FROM sensors [FROM NOW-5
MINUTES TO NOW]
WHERE prox >10
```

GSN

- Sensor middleware
- REST API

```
http://montblanc.slf.ch:22001/multidata?vs[0]=sens
ors&
field[0]=proximity_field&c_min[0]=10&
from=15/05/2012+05:00:00&to=15/05/2012+10:00:
00
```

Cosm/Xively

- Sensor middleware
- Open platform
- REST API

```
http://api.cosm.com/v2/feeds/14321/datastreams/
4?start=2012-05-15T05:00:00Z&end=2012-05-
15T10:00:00Z
```

Underlying Query Processors

SPARQLStream

```
SELECT ?proximity
FROM STREAM <http://streamreasoning.org/SensorReadings.srdf>
[NOW-5 S]
WHERE {
  ?obs a ssn:ObservationValue;
  qudt:numericalValue ?proximity;
  FILTER (?proximity>10) }
```

R2RML

Query
rewriting

$\pi_{\text{timed, prox}}$
 $\sigma_{\text{prox}>10}$
 $\omega_{5 \text{ Seconds}}$
sensors

```
SELECT prox FROM sensors [FROM NOW-5 MINUTES TO NOW]
WHERE prox >10
```

SNEE (DSMS)

```
SELECT prox FROM sensors.win:time(5 minute)
WHERE prox >10
```

Esper (CEP)

```
http://montblanc.slf.ch:22001/multidata?vs[0]=sensors&field[0]=proximi
ty_field&c_min[0]=10&
from=15/05/2012+05:00:00&to=15/05/2012+10:00:00
```

GSN (mddlwr)

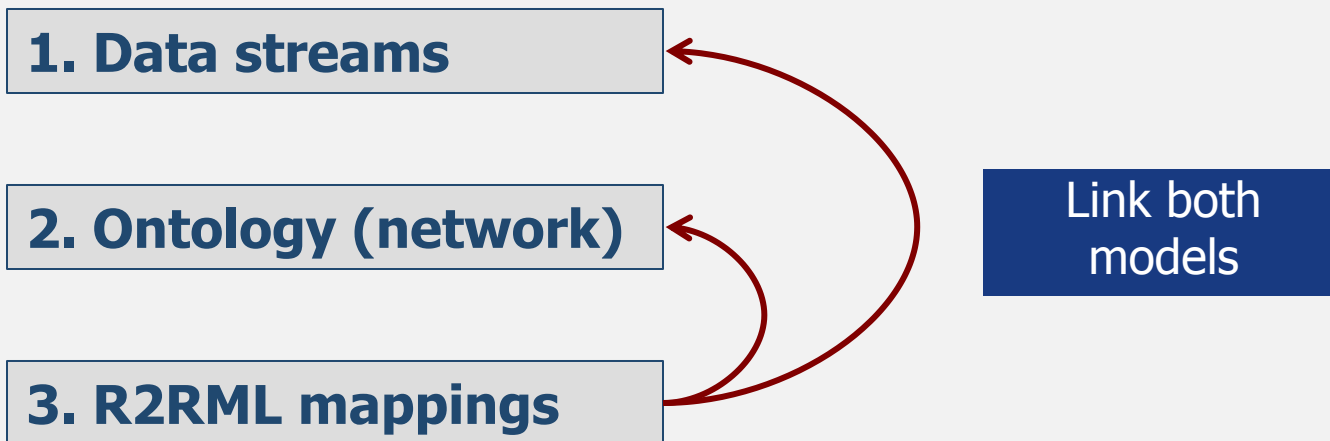
```
http://api.cosm.com/v2/feeds/14321/datastreams/4?start=2012-05-
15T05:00:00Z&end=2012-05-15T10:00:00Z
```

Cosm Xively

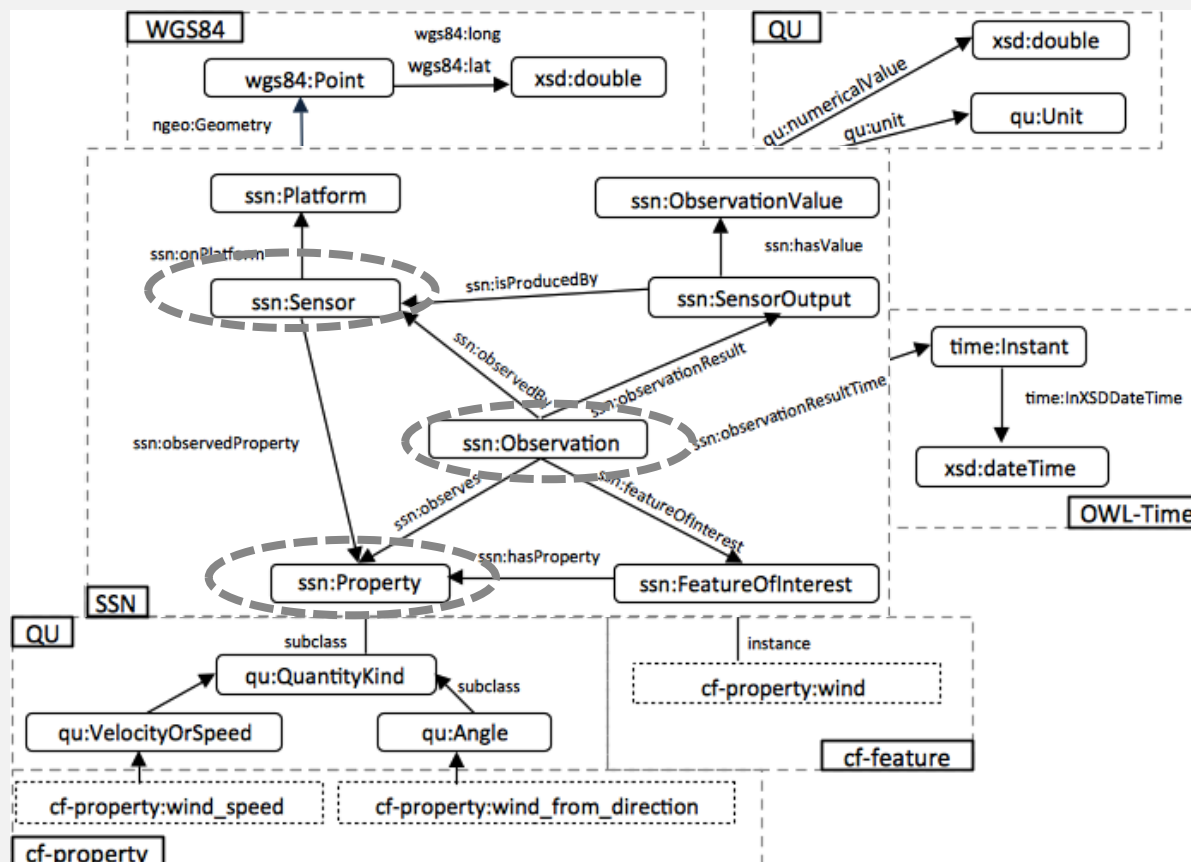
Underlying query processors

Features	Esper	SNEE	GSN	Cosm/Xively
Projection	✓	✓	✓	Fixed
Proj expression	✓	✓	✗	✗
Joins	✓	✓ ✗ only window	✗	✗
Union	✗	✓ ✗ not windows	✓	✗
Selection	✓	✓	✓	✗ ✓ limited
Aggregates	✓	✓	✓ ✗	✗
Time window	✓	✓	✓	✓
Tuple window	✓	✓	✓	✗
R2S	✓	✓	✗	✗
Conjunction, Disj	✓	✗	✗	✗
Repetition pattern	✓	✗	✗	✗
Sequence	✓	✗	✗	✗

- Main ingredients:



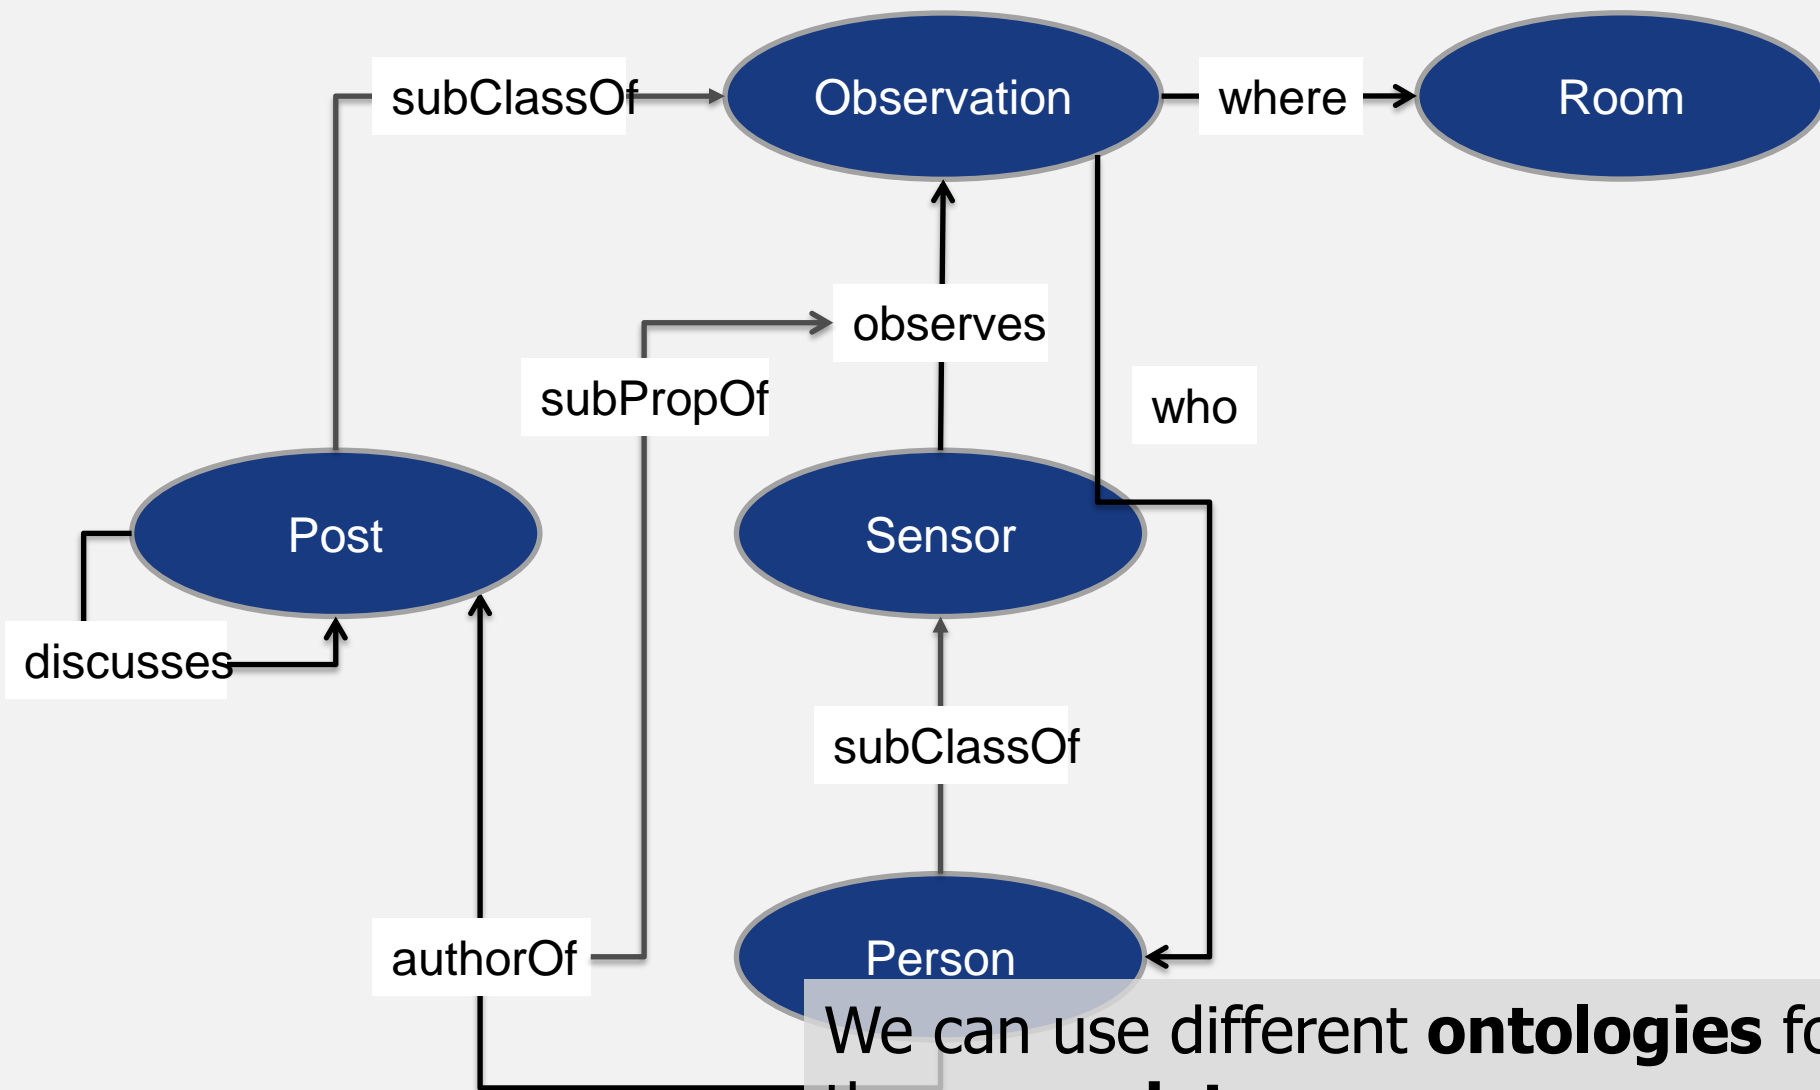
W3C SSN Ontology



modeling our streaming data

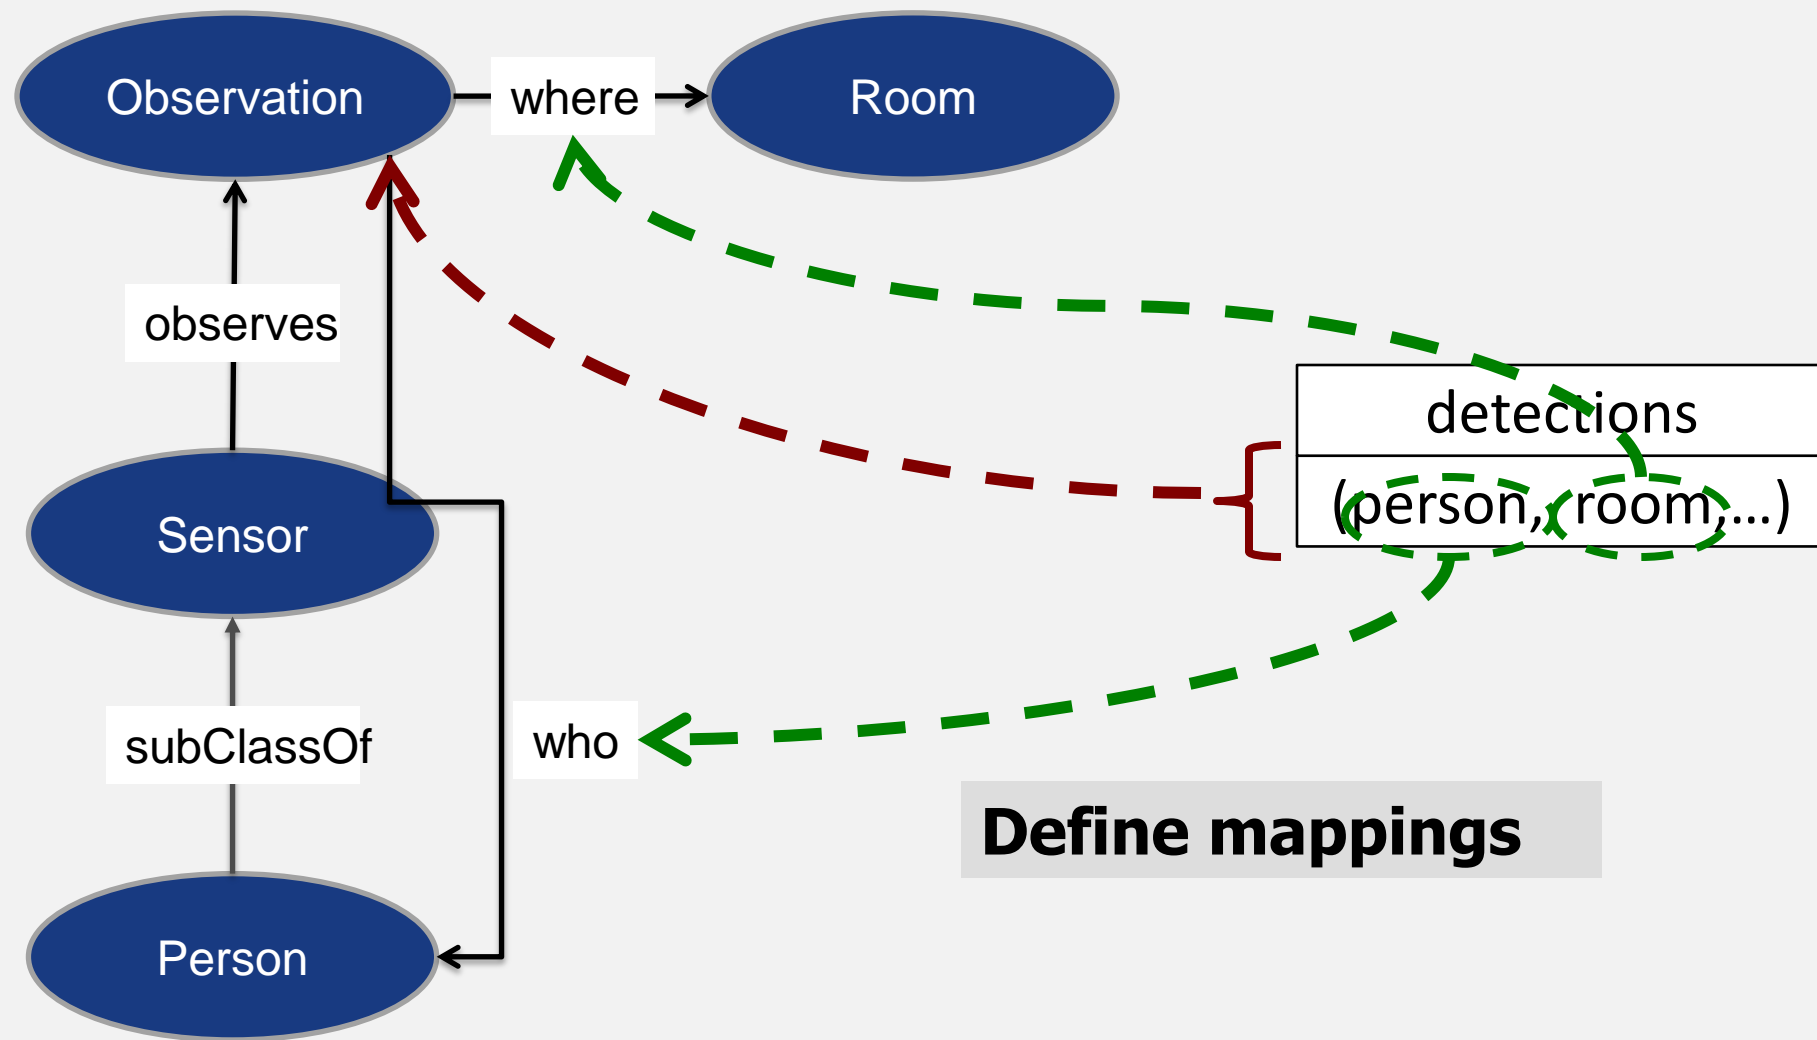
combine with **domain ontologies**

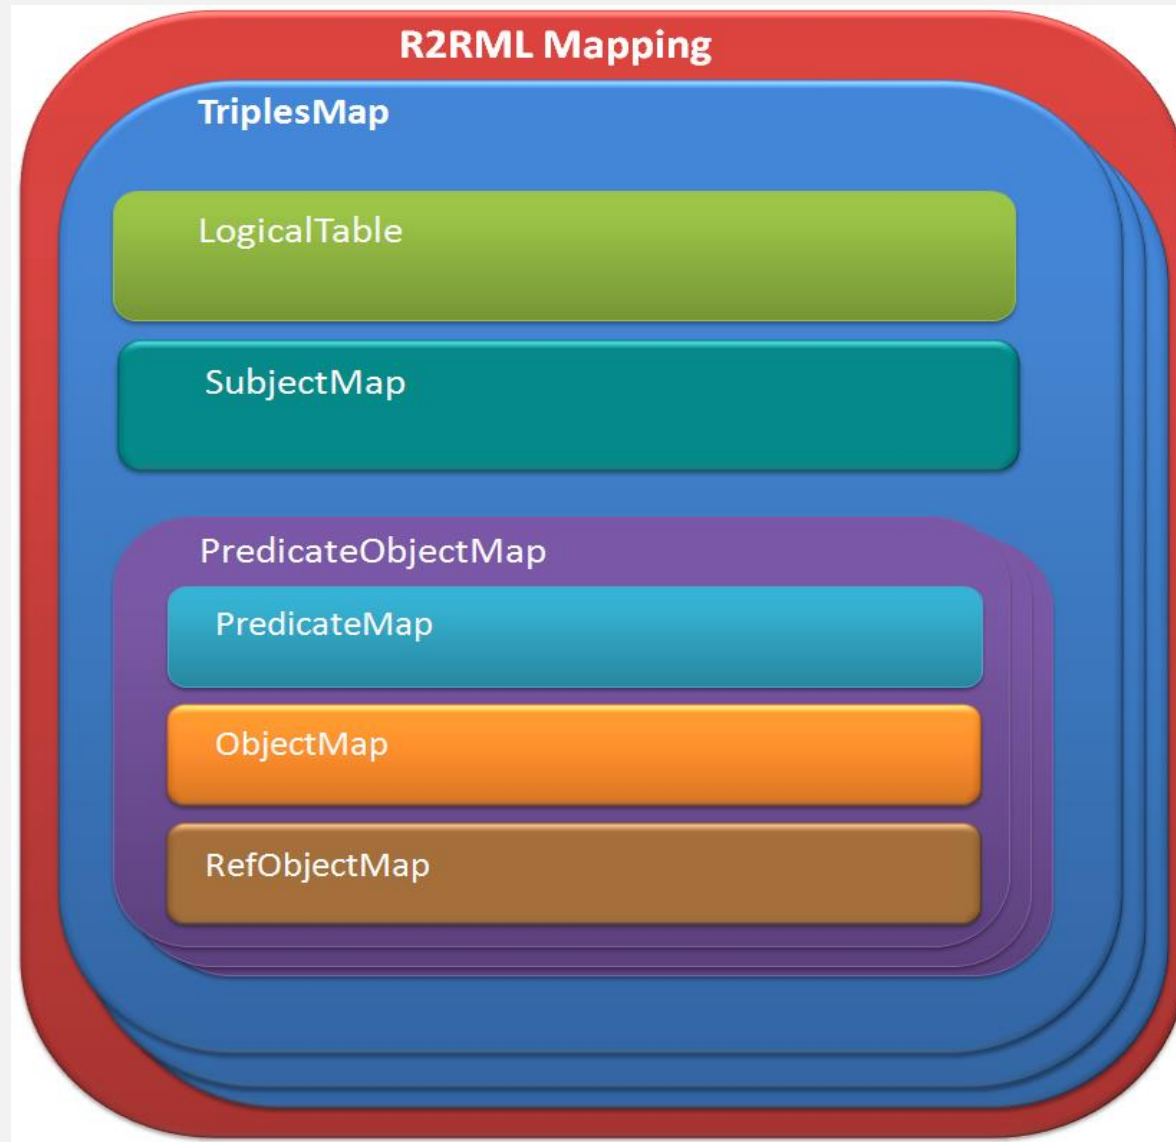
Our simpler ontology...

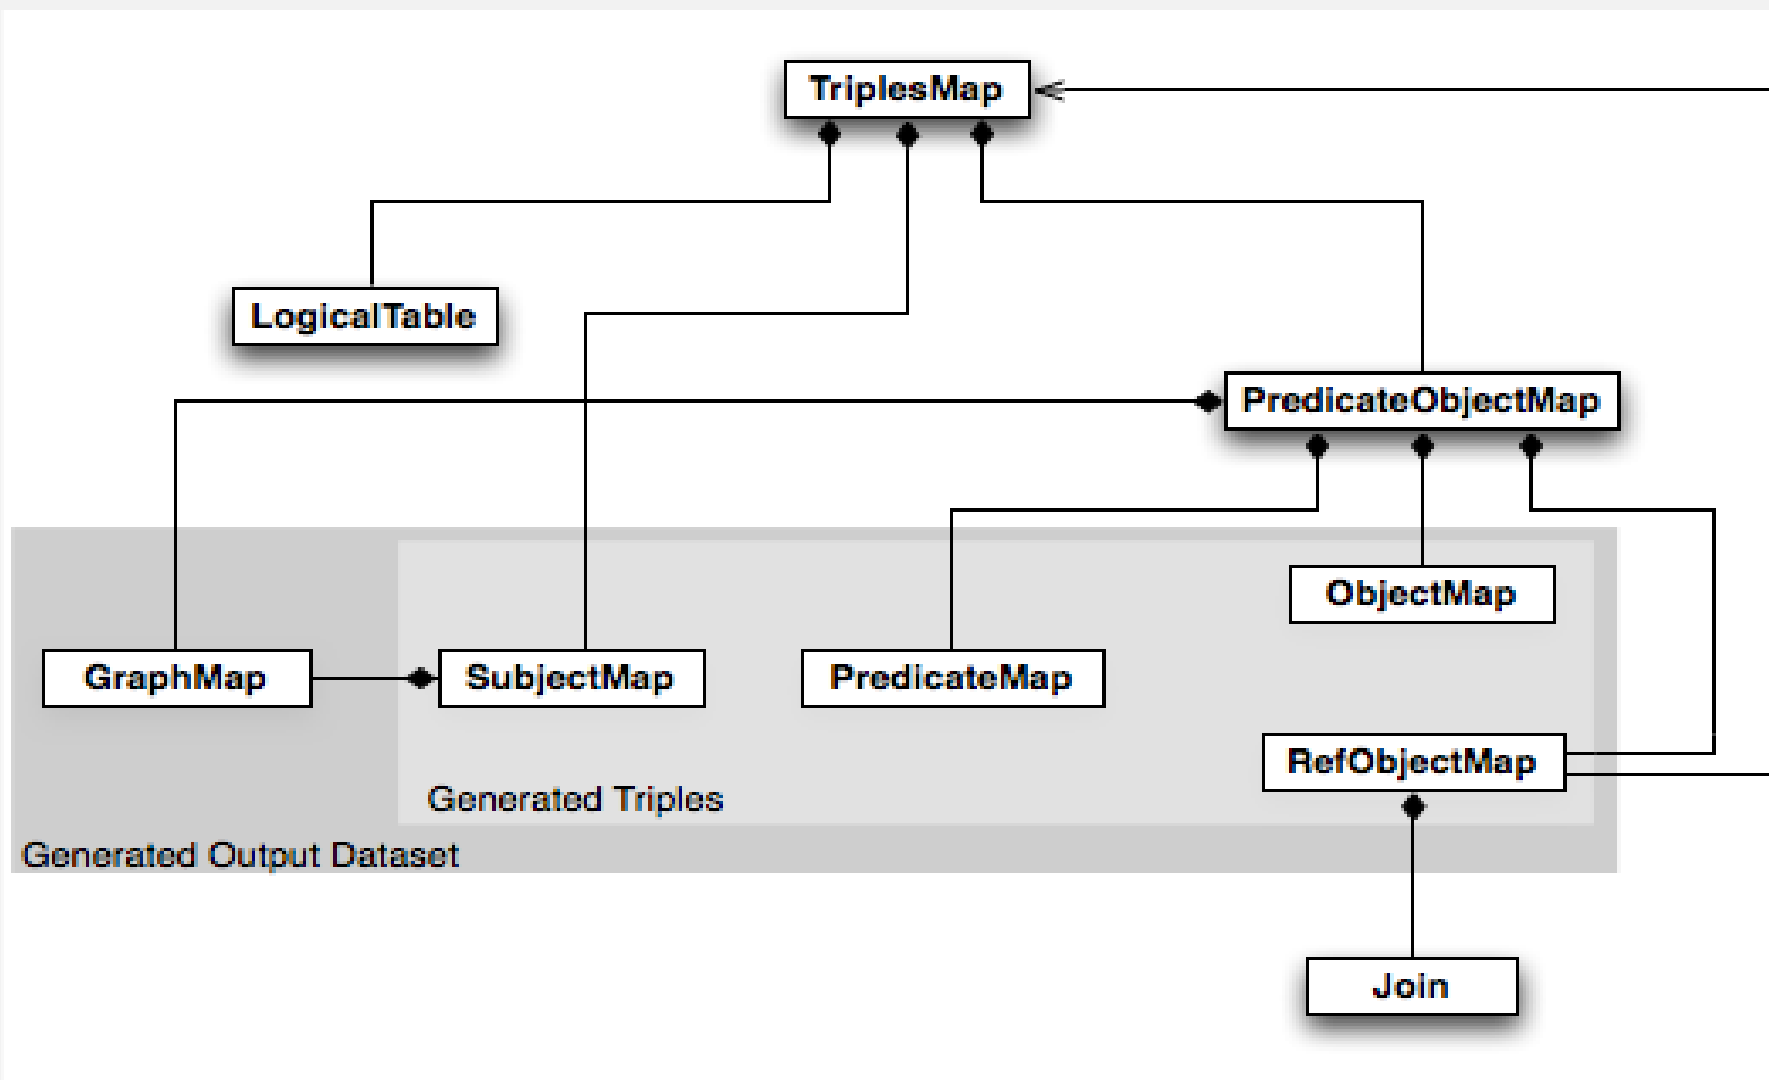


We can use different **ontologies** for the **same data**

Our simpler ontology...







Mapping definition

```
:triplesMap a rr:TriplesMap;  
  rr:logicalTable [ rr:tableName "sensors"; ]  
  
  rr:subjectMap [  
    rr:template "http://streamreasoning.org/data/Observation/{person}{timed";  
    rr:class sr4Id:Observation; rr:graph sr4Id:socialstream.srdf ];  
  
    rr:predicateObjectMap [  
      rr:predicate sr4Id:who ;  
      rr:objectMap [ rr:template "http://streamreasoning.org/data/Person/{person" ]];
```

the stream
name

stream
attributes

subject URI

triple predicate + object

the object (a URI in this case)

Morph-streams:

- Coded in **Scala**
- JAR bundle, use it from Scala or **Java** code
- Maven, Sbt
- **Examples**
 - **One off query**
 - **Register continuous query**
 - **Pull data**
 - **Push**
 - **Basic REST**
- <https://github.com/jpcik/morph-streams>

- Parse SPARQLStream

```
val query= "PREFIX sr4ld: <...>. SELECT ?a ..."  
val syntax= StreamQueryFactory.create(query);
```

- Execute One-off query

```
val query= "PREFIX sr4ld: <...>. SELECT ?a ..."  
mapping=Mapping(new URI(mappings/social.ttl))  
val adapter:QueryEvaluator=Application.adapter(system)  
val results= adapter.executeQuery(query,mapping)
```

Mapping

Bindings

■ Register and Pull

```
val queryid= adapter.registerQuery(query,mapping)
val results1=adapter.pull(queryid)*
val results2=adapter.pull(queryid)
```

Query identifier

■ Register and Push

```
class ExampleReceiver extends StreamReceiver{
  override def receiveData(s:SparqlResults):Unit=
    Logger.debug("got: "+res)
}
val receiver=new ExampleReceiver
val queryid= adapter.listenToQuery(query,mapping,receiver)
```

Implement receiver

- `encoded_value=$(python -c "import urllib; print urllib.quote('\"SELECT
DISTINCT ?timeto ?obs FROM NAMED STREAM
<http://emt.linkeddata.es/data#busstops.srdf> [NOW - 30 S] WHERE
{ ?obs a <http://emt.linkeddata.es/data#BusObservation>. ?obs
<http://purl.oclc.org/NET/ssnx/ssn#observedBy> <http://transporte.li
nkeddata.es/emt/busstop/id/2018>. ?obs
<http://purl.oclc.org/NET/ssnx/ssn#observationResult> ?output.
?output <http://emt.linkeddata.es/data#timeToBusValue> ?av. ?av
<http://data.nasa.gov/qudt/owl/qudt#numericValue> ?timeto. }\"'))")`

Just encoding query

- `curl
"http://streams.linkeddata.es/emt/sparqlstream?query=$encoded_val
ue"`

Disclaimer: Simplistic, not implementing all of the SPARQL protocol

```
{
  "head": {
    "vars": [ "timeto" , "obs" ]
  },
  "results": {
    "bindings": [
      {
        "timeto": { "datatype": "http://www.w3.org/2001/XMLSchema#string" , "type": "typed-
literal" , "value": "0" } ,
        "obs": { "type": "uri" , "value":
"http://transporte.linkeddata.es/emt/busstop/id/2018/busline/9/observation/20/09/2013%2010:2
8:19%20%2B0200" }
      }
    ]
  }
}
```

Bindings in JSON



- Morph-Streams
 - <https://github.com/jpcik/morph-streams>
- See demos
 - <http://transporte.linkeddata.es/> (SPARQL-Stream tab)
 - Check our Madrid buses demo at SSN2013 workshop tomorrow
- Read out more
 - Enabling Query Technologies for the Semantic Sensor Web.
J.-P. Calbimonte, H. Jeung, O. Corcho and K. Aberer.
International Journal on Semantic Web and Information Systems
IJSWIS, Volume 8(1)., 2012
- Contact point
 - jp.calbimonte@upm.es
 - ocorcho@fi.upm.es

Tutorial on RDF Stream Processing

M. Balduini, J-P Calbimonte, O. Corcho,
D. Dell'Aglio, E. Della Valle

<http://streamreasoning.org/rsp2014>



SPARQLStream: Ontology- based access to data streams

Jean Paul Calbimonte, Oscar Corcho

jp.calbimonte@upm.es, ocorcho@fi.upm.es

<http://www.oeg-upm.net/>



$\langle s, p, o \rangle$



$\langle \text{aemet:observation1}, \text{ssn:observedBy}, \text{aemet:Sensor3} \rangle$

$\langle \text{aemet:observation1}, \text{qudt:hasNumericValue}, "15.5" \rangle$

For streams?

$(\langle s, p, o \rangle, \tau)$

$(\langle \text{aemet:observation1}, \text{qudt:hasNumericValue}, "15.5" \rangle, 34532)$

timestamped triples

- Gutierrez et al. (2007) Introducing time into RDF. IEEE TKDE
- Rodríguez et al. (2009) Semantic management of streaming data. SSN